


MODUL UNIVERSITY VIENNA

SESSION 7

Frameworks, Risk, and the EU AI Act

David B. Smith



Instructor Notes: Frame this as a structural shift in the course:

- Prior sessions focused on how individuals reason ethically (interpretation, judgment, perspective)
- This session examines what changes when reasoning must be made consistent, repeatable, and transferable across individuals

Emphasize that this is not replacing ethical reasoning, but placing it under constraint.

Intellectual throughline:

- Ethical reasoning produces varied but defensible answers
- Real-world systems require consistent and enforceable decisions
- This tension motivates the need for frameworks

This is the beginning of a move from open-ended reasoning
→ structured systems of evaluation.

From Frameworks to Regulation

SEID 2364 — Week 7

Instructor Notes: Extend Slide 1 into a clearer argument:

- Ethical reasoning allows disagreement while still being valid
- Frameworks reduce disagreement by constraining how decisions are made

Key pivot:

- Multiple frameworks can exist, each grounded in different underlying principles
- Regulation is the enforced selection of a specific framework for shared use

Key idea:

- The question is no longer "what do you think?"
- It becomes "which framework has been selected—and how is it enforced?"

Do not open discussion yet—this is conceptual orientation.

SLIDE 2A — SYSTEM DESCRIPTION

Personalized Content Ranking and Moderation

(e.g., YouTube / TikTok)

- Analyzes user behavior (views, likes, watch time, interactions)
- Ranks and selects content for each individual user
- May remove, suppress, or promote content
- Produces a personalized feed per user (no single global output)

Instructor Notes: Clarify the system precisely:

- There is no single "output"—each user experiences a different version
- The system is adaptive and continuously updating

Why this system:

- It is familiar
- It produces both clear benefits and contested harms
- It contains hidden processes (ideal for later slides)

Do not allow evaluation yet—this is definition only.

Provocation 1

List the potential benefits and harms this system can produce.

- Be sure to consider that there are multiple types of harms/benefits
- These may differ across stakeholders (individual users, creators, organizations, and the platform).

Instructor Notes:

Purpose:

- Surface intuitive judgments
- Establish baseline variation in reasoning

Expect:

- Broad, unsystematic answers
- Mixing of harms, benefits, and assumptions

Do not correct yet—this variability is necessary for the next steps.

Provocation 2

Describe the conditions under which this system's harms are likely to outweigh its benefits.

(features of the platform, user context, and feedback dynamics)

Instructor Notes: Shift the mode of reasoning:

- from opinion → conditional analysis

Push students to:

- identify when harm emerges
- connect harm to system behavior (not abstract claims)

Intellectual move:

- Harm is not inherent—it is context-dependent
- Systems produce different outcomes under different conditions

If responses stay vague:

- ask "what has to change for that harm to occur?"

Provide concrete examples if students remain abstract:

- high engagement optimization leading to amplification of extreme or polarizing content → potential harms include:
 - individual users → increased anxiety, distorted perception

of reality, or movement toward extreme belief systems (psychological harm)

- specific groups (e.g., minorities, targeted communities) → increased harassment, hate speech, or coordinated attacks (reputational and social harm)
- politically undecided or vulnerable populations → manipulation of beliefs or decisions (civic harm)
- broader society → erosion of shared informational ground necessary for collective decision-making (societal/civic harm)
- institutions (e.g., elections, public health systems) → reduced trust and effectiveness due to misinformation or polarization (structural harm)
- cross-cutting: erosion of trust (interpersonal, institutional, informational) → reduced cooperation, compliance, and legitimacy across systems

Emphasize explicitly:

- polarization itself is not the harm
- the harm emerges through its effects on specific individuals, groups, or institutions
- recommendation loops reinforcing a narrow range of viewpoints (filter bubbles)
- removal or suppression of legitimate content due to automated moderation errors
- targeted content influencing vulnerable populations (e.g., minors, politically undecided users)
- uneven content exposure across demographic groups (algorithmic bias)
- rapid spread of misinformation during breaking events
- shifts in user behavior (e.g., increased time spent, reduced exposure to diverse perspectives)

Also require parallel articulation of benefits (same structure: behavior → beneficiary → type of benefit):

- high engagement optimization → individual users → relevant, engaging content discovery (utility/experience benefit)
- personalization → niche communities/creators → increased visibility and audience matching (economic/cultural benefit)

- moderation systems → users and public → reduced exposure to harmful or illegal content (safety benefit)
- recommendation diversity tuning → users → exposure to new topics or perspectives (informational benefit)
- rapid information dissemination → public → timely awareness during events (civic/informational benefit)
- platform analytics → creators/deployers → improved content strategy and resource allocation (economic/operational benefit)

Prompts to enforce symmetry:

- "Who benefits, specifically?"
- "What actually improves for them?"
- "Under what conditions does that benefit hold—or break?"

Use these to push conditional framing:

- "Under what system settings or behaviors would this occur?"
- "What would have to be true about the algorithm or data?"

Provocation 3

**What evidence would
be required to support
your claims in
Provocation 2?**

Instructor Notes:

Purpose:

- transition from claims → justification

Clarify explicitly:

- A claim without evidence cannot be evaluated consistently

Push toward specificity:

- behavioral data
- system outputs
- user impact metrics
- moderation decisions

Provide concrete examples if students remain abstract:

- engagement statistics showing shifts in user behavior (e.g., increased watch time on specific content types)
- comparative feed outputs across different users (to detect personalization effects or bias)

- records of content amplification or suppression (what was boosted vs hidden)
- moderation logs showing what content was removed or flagged and why
- reports of user outcomes (e.g., exposure to misinformation, behavioral changes)
- experimental or A/B testing results showing how changes to the algorithm alter outcomes
- external studies or journalism documenting platform effects on groups or public discourse

Intellectual move:

- ethical reasoning → epistemic requirement (what must be known)

Provocation 4

**Where in the system
does that evidence
exist?**

Instructor Notes:

Purpose:

- situate evidence within a structure

Move students to think in terms of systems:

- data collection layer
- model processing
- output/interface layer

Key idea:

- Evidence is not abstract—it is located and mediated

Provide concrete examples if students remain abstract:

- user interaction logs (watch time, clicks, engagement data)
- recommendation outputs (what content is shown to whom)
- moderation records (removed, flagged, or down-ranked content)

- training data sets or sampled datasets
- internal model metrics (accuracy, bias measurements)
- A/B testing results or system experiments
- external audits or research studies analyzing platform behavior

Bridge to later:

- this becomes SVD (Source → Transformation → Destination)

Provocation 5

Who owns, controls, or has access to that evidence?

Instructor Notes:

Purpose:

- introduce governance and asymmetry

Clarify dimensions:

- ownership → legal or institutional control
- control → ability to modify or filter
- access → ability to view or retrieve

Critical insight:

- Evidence may exist but be unavailable
- Different stakeholders have unequal access

This is the first introduction of power into the analysis.

What criteria did you use to make these decisions?

SEID 2364 — Week 7

Instructor Notes:

Purpose:

- make implicit reasoning explicit

Students should identify:

- fairness
- outcomes
- autonomy
- intent

Do not resolve differences yet.

Key move:

- show that different criteria produce different conclusions.

SLIDE 9 — EMERGING DISTINCTION

Ethical Reasoning vs Frameworks

Ethical Reasoning

- Context-dependent
- Allows ambiguity
- Tolerates disagreement
- Variability across actors

Frameworks

- Constrained interpretation
- Designed for repeatability
- Same inputs → same outputs
- Consistency across actors

Instructor Notes:

Purpose:

- Synthesize what just happened in Slides 3–8
- Make visible that different students reached different conclusions using different criteria

Up to this point, we have been working within ethical reasoning systems—frameworks that allow for interpretation, judgment, and disagreement. As we have seen, different ethical frameworks can produce different, yet internally consistent, conclusions about the same system. These differences emerge because each framework relies on distinct underlying principles—axiom-like assumptions about what is considered "good," "fair," or "just."

At the individual level, we can accept this variability.

Shared systems exist at multiple levels of society

(individual, institutional, societal),

Each level may deploy different underlying philosophical approaches.

Even at the same level, different institutions (e.g., a business and a school) may adopt distinct or conflicting ethical frameworks based on their goals, incentives, and responsibilities.

This creates additional complexity: systems that interact across levels may display conflicting ethical conclusions, even before formal evaluation begins. One can adopt and justify a preferred framework, and disagreement can remain productive.

However, once we move from individual judgment to shared systems—systems that affect many people, operate at scale, or require oversight—this variability becomes a problem. If different evaluators apply different criteria, the same system may be judged inconsistently, leading to instability, lack of accountability, contested outcomes, and erosion of trust—reducing compliance, weakening coordination, and undermining institutional legitimacy.

This introduces a constraint: to be consistent, we require consistent criteria.

To achieve this, we move from open-ended ethical reasoning to the construction of standards—frameworks that define categories, constrain interpretation, and enable repeatable evaluation across different actors.

The transition, then, is not from "ethics" to "regulation," but from interpretive evaluation to formalized, enforceable classification.

Optional extension: Note that consistency introduces rigidity—this tradeoff will be examined later.

SLIDE 10 — FRAMEWORKS

Frameworks

Defined criteria

- Frameworks specify what is important

Constrained interpretation

- Frameworks constrain how judgments are made

Repeatable evaluation

- Frameworks enable repeatability across different actors

Instructor Notes: Define frameworks as intermediate structures between open ethical reasoning and regulation:

- frameworks specify what counts (criteria)
- frameworks constrain how judgments are made
- frameworks enable repeatability across different actors

Clarify:

- multiple frameworks can exist simultaneously
- each is grounded in different underlying (axiom-like) principles

Key point:

- frameworks reduce variability, but do not eliminate disagreement
- different frameworks can still produce different conclusions

Bridge explicitly:

- the problem is not just reasoning—it is which framework is

used

- context-dependent
- allows ambiguity
- tolerates disagreement

Important clarification:

- Ethical reasoning is not "wrong" or "insufficient"
- It is simply not designed for consistency across actors

Prompt if needed:

- "Could two reasonable people disagree and both be justified?" → yes

SLIDE 11 — REGULATION

Regulation (Enforcement)

Selection

- Regulation selects one framework (or a defined set) for shared use

Enforcement

- Regulation enforces that selection across actors and contexts

Standardization

- Same system should yield the same classification across evaluators

Instructor Notes: Define regulation precisely:

- regulation selects one framework (or a defined set) for shared use
- regulation enforces that selection across actors and contexts

Core idea:

- regulation = enforced selection of frameworks

Contrast with Slide 10:

- frameworks → multiple options can coexist
- regulation → limits which frameworks are allowed in practice

Implications:

- same system should yield the same classification across evaluators
- disagreement is reduced by removing alternative frameworks

Introduce consequence:

- consistency is achieved through constraint
- this may exclude alternative, but still valid, ethical perspectives

Bridge forward:

- the EU AI Act is an example of this process in action
- systems that constrain how decisions are made
- designed to produce the same outcome given the same inputs

Key contrast with Slide 10:

- Ethical reasoning → variability
- Frameworks → repeatability

Introduce tension explicitly:

- frameworks gain consistency by reducing nuance

SLIDE 12 — PRESSURE QUESTION

What must be true for consistent evaluation to work?

Instructor Notes: This is a synthesis question—students should now draw from earlier slides.

Guide toward three necessary conditions:

1. shared criteria (Slide 8)
2. access to evidence (Slides 5–7)
3. clear system specification (implicit but emerging)

If needed, push:

- "Could you apply your reasoning if someone else used different criteria?"
- "Could you apply it without access to evidence?"

Consistency requires specification

SEID 2364 — Week 7

Instructor Notes: State this as a formal conclusion:

- Without specifying the system, the criteria, and the evidence, consistent evaluation is impossible

Refine "specification" using an engineering analogy:

- Think of a specification sheet: it defines the system through clear, structured, and often quantitative descriptors
- It enables different engineers (or evaluators) to interpret the same system in the same way

Clarify explicitly:

- specification here means operational, structured, and interpretable definition
- not vague or arbitrary instruction

Important distinction:

- quantification may be part of specification (metrics, thresholds, measurements)
- but only when embedded within a clearly defined

structure

Make explicit connection:

- what you did in Slides 4–7 was moving toward this kind of specification

This is the pivot from:

- exploratory reasoning → formal systems



Instructor Notes: Frame the EU AI Act not as law first, but as a design response to the problem just identified.

Key framing:

- The EU AI Act attempts to:
 - define categories
 - enforce consistency
 - require documentation and justification

Bridge statement:

- "What we just struggled with informally is exactly what regulatory frameworks attempt to solve formally"

SLIDE 14A — ORIGINS

Origins: HLEG (2018–2020)



Motivated by growing concerns about AI systems at scale

- Concerns: bias, opacity, manipulation, and lack of accountability
- Tasked by the European Commission to translate ethics into actionable guidance
- Produced Ethics Guidelines for Trustworthy AI (2019)
- Introduced principles: human oversight, transparency, accountability, risk-based thinking

Instructor Notes: Position the HLEG as a compression layer between broad ethical discourse and formal regulation:

- many perspectives → distilled into structured principles
- not a single author; a curated, multi-stakeholder expert group (~52 members)

Clarify the emergent capabilities that motivated concern (circa 2016–2019):

- large-scale recommendation systems shaping behavior and information exposure (YouTube, Facebook feeds)
- early deepfakes and synthetic media (GAN-based image/video manipulation)
- advances in computer vision (facial recognition in public and commercial systems)
- natural language systems enabling automated content generation and moderation
- predictive systems used in hiring, credit scoring, and policing (algorithmic decision-making)
- large-scale data aggregation enabling behavioral profiling

and targeted influence (e.g., Cambridge Analytica)

Key insight:

- these systems were not just tools—they were beginning to mediate perception, decision-making, and social structures at scale

Key move:

- converts implicit ethical assumptions into explicit, shareable criteria

Bridge forward:

- these principles inform later regulatory specification
→ the EU AI Act

SLIDE 15

EU AI Act: Context and Authority



Why it was created, when it emerged, and what power it holds

- Created to address risks of AI systems at societal scale (safety, rights, trust)
- Proposed by the European Commission (2021)
- Formally adopted by the European Union (2024)
- Applies across EU member states
- Has legal force: organizations must comply or face penalties

Instructor Notes: Frame the EU AI Act as:

- a response to the exact problem identified in prior slides: inconsistent evaluation across actors
- an attempt to create a shared, enforceable specification and classification system

Clarify motivation:

- protect fundamental rights (privacy, fairness, non-discrimination)
- ensure safety and accountability in AI deployment
- maintain public trust in large-scale systems

Clarify power:

- applies to any system deployed within or affecting the EU market
- includes enforcement mechanisms (fines, restrictions, bans)
- shapes how systems are designed, not just evaluated

Key conceptual link:

- this is not just policy—it is a regulatory encoding of a selected framework



Instructor Notes:

Keep presentation minimal, but sharpen the conceptual function:

Key idea:

- complex systems are reduced to discrete, mutually exclusive categories

Clarify purpose of classification:

- determines obligations
- determines oversight
- determines whether the system is allowed at all

Add explicit constraint:

- classification must be applied consistently across systems and evaluators

Introduce upcoming structure (important for 16A–D):

- each category corresponds to a different regulatory

response

- prohibition
- strict control
- transparency
- minimal intervention

Framing line (use verbally):

"The system does not change—only how it is categorized and regulated."

Prompt (optional):

- "Could two regulators classify the same system differently?"
- this is the problem the framework is designed to reduce

SLIDE 16A — UNACCEPTABLE RISK

Banned Systems (Prohibited)



Systems that fundamentally violate rights or autonomy and are therefore not allowed to be deployed.

- Manipulation that undermines autonomy
- Exploitation of vulnerable groups
- Social scoring
- Real-time biometric surveillance (limited exceptions)

Prohibited

Unacceptable Risk

Instructor Notes:

Operational definition:

- systems fundamentally incompatible with rights and dignity

Key point:

- not regulated → prohibited

Emphasize:

- this is a normative boundary, not a technical threshold
- the system is removed entirely from the space of acceptable solutions

SLIDE 16B**Allowed with Strict Obligations**

Systems that significantly affect safety or rights and are permitted only with strict oversight, documentation, and control.

- Hiring and recruitment systems
- Credit scoring / loan decisions
- Education (grading, admissions)
- Medical AI
- Critical infrastructure

High Risk

Strict Obligations

Instructor Notes:

Operational definition:

- systems that significantly affect safety, rights, or access

Core requirements (conceptual, not exhaustive):

- risk management
- data governance
- documentation and traceability
- human oversight

Key point:

- most contested real-world systems will fall here
- regulation here does not eliminate the system—it constrains it

SLIDE 16C — LIMITED RISK

Transparency Required



Systems that may influence or mislead users and therefore require disclosure but not restriction.

- Chatbots
- AI-generated media (deepfakes)
- Systems that simulate human interaction

Limited Risk

Transparency
Required

Instructor Notes:

Operational definition:

- systems that may mislead users if not disclosed

Regulatory response:

- disclosure, not restriction

Key point:

- concern is awareness, not direct harm
- the user must know they are interacting with or seeing AI

SLIDE 16D — MINIMAL RISK

No Specific Regulation



Systems that pose minimal risk and are largely unregulated under the Act.

- Recommendation systems
- Spam filters
- Game AI
- Basic personalization

Minimal Risk

No Specific Regulation

Instructor Notes:

Operational definition:

- negligible impact on rights or safety

Key point:

- majority of systems fall here
- classification is not always obvious—context can shift a system upward

Important tension:

- systems may appear minimal risk but produce cumulative or indirect harms

SLIDE 17 — APPLICATION



How can the EU AI Act's risk classification be applied to the personalized content ranking system?

Use the harms you identified to determine how this system maps onto EU AI Act risk levels.

Instructor Notes:

Purpose:

- Transition from understanding → classification using prior analysis
- Students are not generating new harms—they are interpreting existing ones

Set the role explicitly:

- Students are acting as regulators applying a fixed framework

Task:

- assign the system to one category only
- justify using:
 - harms identified (Slides 4–7)
 - how those harms are produced
 - who is affected

Key clarification:

- the question is not "how much harm exists"
- the question is:
 - how directly the system produces those harms
 - whether those harms affect rights, safety, or access

Guiding questions:

- "Do these harms emerge through influence, or direct decision-making?"
- "Do they affect rights or access in a structured way?"
- "Are the harms inherent, or conditional?"

Important constraints:

- criteria are partially specified
- system boundaries are ambiguous
- evidence is incomplete

Key dynamic:

- students must classify despite uncertainty

Expected outcome:

- disagreement
- reframing of the system
- requests for missing information

Do not resolve too quickly—this tension is the point.

Framing line:

"You are not evaluating the system—you are assigning it a regulatory status based on the harms you identified."

SLIDE 18 — INFORMATION GAP

What information do we need to make this classification?

Instructor Notes: Reconnect to Slides 5–7.
Students should recognize missing evidence and access.

SLIDE 19 — MISSING INFORMATION

Is that information accessible? Why or why not?

WE MADE IT TO HERE 2026-04-16

Instructor Notes: Highlight constraint:

- decisions must be made without full knowledge

SLIDE 20

Why Is Classification Difficult?

- What information is missing?
- Where does that information exist?
- Who has access to it?

Instructor Notes: This slide remains focused on the shared system and prepares the class for mediation analysis.

Purpose:

- make explicit why the EU AI Act is difficult to apply cleanly
- surface the limits of classification when systems are partially visible

Clarify:

- this is not yet about their own systems
- it is a guided reflection on the shared example

Framing line: "Before moving to your own systems, we need to understand why this shared system is difficult to classify."

Time: 5–7 minutes discussion

SLIDE 21 — BBS INTRODUCTION

Mediation Model



Instructor Notes: Introduce BBS mediation here as a tool for explaining the difficulty identified in Slide 21.

Purpose:

- show that information is not directly available
- show that evidence is produced, transformed, and distributed through pathways

Connect explicitly to EU framing:

- classification depends on evidence
- evidence depends on where it sits in the mediation pathway
- therefore mediation analysis helps explain why regulatory classification is difficult

Use the shared system:

- identify where data originates
- what transforms it
- what outputs are produced

Framing line: "We now need a structural model of the system so we can understand why EU risk classification is difficult."

SLIDE 22 — YOUR SYSTEM

Journal Reflection — Apply the Same Structure to YOUR System

- What is your system trying to do?
- Where does its data originate?
- What transforms that data?
- What outputs or decisions does it produce?

Instructor Notes: This is the explicit pivot to the students' own systems.

State clearly:

- up to this point, the class has used the shared content-ranking system
- now students transfer the same analytical structure to their own systems

Purpose:

- begin assignment-oriented system analysis
- use mediation pathways to inform later EU AI Act classification

Framing line: "We now move from the shared example to your own system. Use the same mediation structure to begin describing it."

Constraint:

- this is still exploratory
- no final classification yet

Time: 8–10 minutes journal writing

SLIDE 23 — VISIBILITY

Where is required information located?

Instructor Notes: Clarify that this question is about location of evidence, not interpretation:

- Where in the system would the information needed for evaluation actually reside?
- Examples: logs, outputs, training data, moderation records, internal metrics

Key distinction:

- information is embedded within system components (data layer, model layer, interface layer)
- it is not abstract or freely available

Important conceptual move:

- students often assume that if information exists, it is accessible
- this slide separates existence from availability

Bridge forward:

- identifying location sets up the next question: who can

actually access that information?

SLIDE 24 — ACCESS

Who has access to that information?

Instructor Notes: Clarify explicitly what "it" refers to:

- the information required to evaluate and classify the system

Reframe the question:

- which actors have access to that information, and at what level?

Situate the class within this structure:

- students are operating as external evaluators with partial access
- they do not have full visibility into the system

Differentiate actors:

- users → see outputs only (their feed)
- researchers/public → partial, indirect visibility
- deployers/platform → full internal access
- regulators → conditional or negotiated access

Key insight:

- location $\neq$ access
- the same system will appear differently depending on who is observing it

Reinforce connection to Slide 7:

- ownership, control, and access determine what can be known

Bridge forward:

- limited access constrains the ability to justify classification

SLIDE 25 — ASSIGNMENT BRIDGE

From Classification → Justification

Can you justify your classification using:

- System behavior
- Evidence and access
- EU AI Act categories

Instructor Notes: This is not a repeat of earlier classification work.

Clarify the shift explicitly:

- earlier: students classified a shared system (guided, in-class)
- now: students prepare to justify classification of their own system (independent work)

This slide functions as a bridge to the assignment.

Guide briefly as a whole-class framing:

- What counts as a valid justification?
- What evidence would be required?
- Where are the limits of what you can know?

Reinforce dependencies:

- justification depends on Slides 27–28 (visibility and access)

Framing line: "Classification is not the end—the task is to justify it under constraint."

Do NOT have students complete full analysis here.

This sets up Assignment 7:

- formal EU AI Act classification
- structured justification
- application to their system

Time: 5–10 minutes (discussion only)

- this guided process is what students will perform independently on their own systems

SLIDE 28 — REFLECTION

**After applying the EU classification,
what changed in your evaluation?**

Instructor Notes: Compare intuitive vs structured reasoning.

Next: Applying the EU AI Act to Your Own System

Prepare for Assignment 7

SEID 2364 — Week 7

Instructor Notes: Prepare students for Assignment 7.

Emphasize:

- deeper analysis
- independence

SLIDE 26 — EVALUATION FRAME (KEEP ON SCREEN)

Use While Working on YOUR System

Stakeholders

Who is affected
(users, subjects,
deployers, organizations,
public)

Risk Distribution

How impact differs
across stakeholders

Visibility vs Access

Where evidence exists
vs. who can see it

What risks are visible vs legible?

Purpose:

- provide a compact, persistent reference while students work
- align directly with EU AI Act justification needs

Keep this slide visible during work time.

Do not expand—this is a framing tool, not a lecture.

Time: remains on screen during individual work

Clarify distinction:

- visible = observable
- legible = measurable and accessible

Regulation replaces choice with obligation

SEID 2364 — Week 7

Instructor Notes: This is the central pivot of the session.

Clarify explicitly:

- earlier: students could select or implicitly rely on different ethical frameworks
- now: the EU AI Act constrains that choice

Key shift:

- you are no longer choosing how to evaluate
- you must evaluate within the regulatory structure provided

State directly:

- the question is no longer "which framework do you prefer?"
- it is "how does this system fit within the required classification system?"

Important implication:

- disagreement does not disappear

- but it must now be expressed within the constraints of the regulation

Framing line: "Regulation does not eliminate judgment—it constrains where and how judgment can occur."

SLIDE 29 — REFLECTION PROMPTS

Reflection

- What was missing?
- What might the framework miss?

Instructor Notes: Encourage critique of framework itself.

Applying frameworks requires visibility, access, and specification

SEID 2364 — Week 7

Instructor Notes: Restate core triad.

ASSIGNMENT 7

Classify Your System Under the EU AI Act

Using the system you selected earlier in the course:

1. Identify the risk category (unacceptable, high, limited, or minimal)
2. Justify your classification using system behavior, evidence, and access
3. Identify what information is visible, what is legible, and what is missing
4. Describe how stakeholder impact varies across affected groups
5. Reflect on what the framework captures and what it might miss

SEID 2364 — Week 7

Instructor Notes: This is the assignment overview slide. Walk through each requirement briefly.

Key points:

- Students must work independently on their own chosen system
- Classification must use the EU AI Act risk categories introduced this session
- Justification requires concrete evidence—not just opinion
- The visibility/legibility distinction from the evaluation frame applies here
- Encourage students to note where the framework falls short—critical reflection is expected

Remind students:

- This builds directly on the guided exercise from today's session
- The depth of analysis should go beyond what was done in class

- Refer back to the evaluation frame slide for structuring their analysis